

WHIR Accumulation Through Lattice-Based Folding

Paul Louis Robert Quesnot
EPFL

June 4, 2026

Abstract

Accumulation is a natural way to avoid verifying many succinct proofs independently, but applying it to hash-based proof systems requires accumulating verifier claims rather than proof byte strings. We study this question for WHIR, a fast Reed–Solomon proximity-testing based backend, and Symphony-style lattice-based folding.

Our construction models typed accumulation inputs as WHIR proof-validity claims, folds these claims through a lattice-based accumulator, and certifies the folding transition with a commit-and-prove relation whose backend is again WHIR.

The implementation shows that the public-verifier boundary can be preserved in explicit accumulation routes, while identifying the typed CP relation, transcript binding, and relation encoding as the main bottlenecks toward a production-authoritative accumulator.

Keywords. WHIR, accumulation, folding schemes, lattice-based commitments, commit-and-prove SNARKs, interactive oracle proofs, proof systems.

Contents

1	Introduction	4
1.1	Motivation	4
1.2	Problem Statement	4
1.3	Contributions	5
2	Technical Overview	5
2.1	High-Level Architecture	5
2.2	What is Being Accumulated?	6
2.3	Public-Verifier Boundary	6
2.4	Implementation Routes	6
2.5	Performance Direction	7
3	Preliminaries	7
3.1	Interactive Oracle Proofs and SNARKs	8
3.2	WHIR	8
3.3	Accumulation and Folding	9
3.4	Commit-and-Prove SNARKs	9
3.5	Lattice Commitments	9
4	WHIR Accumulation Through Lattice-Based Folding	10
4.1	Interface	10
4.2	From WHIR Proofs to Verification Claims	10
4.3	Lattice-Based Folding of Claims	11
4.4	CP-SNARK for the Folding Transition	12
4.5	WHIR Backend for the CP Relation	13
4.6	Verification Algorithm	14
5	Formal Relation Definitions	15
5.1	WHIR Verification Relation	15
5.2	Folding Relation	15
5.3	Commit-and-Prove Folding Relation	16
5.4	WHIR Backend Relation	17
5.5	Faithfulness of the Encoded Relation	17
5.6	Accumulator Validity and Concrete Route Faithfulness	17
5.7	Row-Level Faithfulness Obligations	18
6	Completeness	20
6.1	Local Completeness of the WHIR Claims	21
6.2	Completeness of Folding	21
6.3	Completeness of the CP Layer	21
6.4	Completeness of the WHIR Backend	22
6.5	Full Completeness Theorem	22
7	Soundness as a Reduction to Row-Level Faithfulness	23
7.1	Assumptions	23
7.2	Soundness Chain	24
7.3	Error Composition	25

8	Implementation	27
8.1	Routes and Data Model	27
8.2	Single-Oracle versus Multi-Oracle Organization	27
8.3	Encoded Folding Transition	28
8.4	Tests and Guardrails	28
9	Evaluation	29
9.1	Experimental Setup	29
9.2	Product Public-Verifier Baseline	30
9.3	K6a Explicit Accumulator Route	31
9.4	N8 Integrated Route and Native Multi-Oracle Evidence	32
9.5	Bottleneck Analysis	34
10	Challenges and Limitations	35
10.1	Relation and Transcript Faithfulness	35
10.2	Prototype Boundaries	35
10.3	Remaining Gaps	35
11	Related Work	36
11.1	WHIR and Reed–Solomon Proximity Testing	36
11.2	Symphony and Lattice-Based Folding	36
11.3	Commit-and-Prove SNARKs	37
11.4	Accumulation and Folding Schemes	37
11.5	Transparent SNARKs and Spartan	37
11.6	Positioning of This Work	37
12	Future Work	38
13	Conclusion	39

1 Introduction

1.1 Motivation

WHIR is already compelling as a proof backend in its own right. It offers a fast hash-based polynomial-commitment architecture, succinct proofs built from Merkle-committed polynomial evaluations, and a public verifier boundary that fits naturally with modern hash-based proof infrastructure. In the current repository, this is not just a design goal: WHIR-backed public verification already exists as an authoritative public-only route. WHIR provides a fast interactive-oracle style polynomial-commitment backend based on Reed–Solomon proximity-testing ideas, and Symphony provides the high-arity lattice-based folding and commit-and-prove setting in which we study accumulation [ACFY24b, Che25].

Even so, verifying many WHIR proofs independently still repeats verifier work. Each proof must be checked against its public inputs, its transcript commitments, its derived challenges, and its folded-output boundary. If an application needs to validate many such proofs, then the backend may already be succinct, but the verifier still performs the same pattern of checks over and over again.

This is exactly the setting in which folding suggests accumulation. Rather than checking many WHIR-backed claims independently, one would like to combine them into a single accumulator transition and certify that transition succinctly. If successful, this preserves the benefits of WHIR as a standalone backend while also amortizing repeated verifier work across many claims.

The subtlety is that the object to be accumulated is not the proof bytes themselves. For WHIR-like systems, what matters is the proof-validity claim induced by each proof: public inputs, Fiat–Shamir commitments, transcript roots/digests, derived challenges, opened evaluations, and folded-output semantics must remain bound to the same verification relation. A sound accumulation scheme must therefore fold verification claims, not raw proof payloads.

1.2 Problem Statement

The mathematical goal of this report is to describe an accumulation procedure for WHIR-backed proof-validity claims. Given an old accumulator Acc_{old} and a batch of WHIR-backed claims, the prover should produce a new accumulator Acc_{new} together with an accumulation proof Π_{acc} such that acceptance implies that the declared accumulator transition is valid and that the accumulated WHIR verification claims hold, except with the explicit error inherited from the underlying backend, folding, and binding assumptions.

To make that goal meaningful, the public-verifier boundary must remain intact. Verification should be public-only and fail-closed. Public inputs, relation metadata, Fiat–Shamir commitments and their roots, challenge derivation, challenge-to- β binding, folded-output algebra, original Ajtai opening validity, original R1CS validity, and accumulator-transition consistency must all be tied to the same proof object. If any of those checks escape into witness-side replay, debug-only helpers, or silent fallback routes, then the resulting object is not a genuine public accumulator.

Within the current implementation, this problem appears through three distinct routes. First, the default public API `verify_public/verify_v2` provides the authoritative WHIR+WHIR public verifier and serves as the baseline public boundary. Second, the explicit opt-in *single-oracle* accumulation route is the current full accumulator workload path and is *NonZK integrity-only*. Third, the repository implements an explicit *multi-oracle* accumulation route for same-shape, nonempty NonZK transitions. These route names are implementation labels; the conceptual distinction is between baseline public verification, explicit full-workload accumulation, and integrated one-proof

accumulation. The single-oracle versus multi-oracle technical distinction is developed in Section 8.2. The rest of the report studies accumulation against this three-route background.

1.3 Contributions

The main contributions of this report are the following.

- **Formalization.** We isolate the right accumulation object for the WHIR setting: proof-validity claims with an explicit public boundary and transcript-binding requirements.
- **Construction.** We describe a WHIR-backed accumulation architecture in which WHIR proof-validity claims are folded through a Symphony-style lattice accumulator and the resulting transition is certified by a CP relation proved with WHIR. We then map this construction to the implemented single-oracle and multi-oracle routes.
- **Completeness and soundness reduction.** We prove completeness of the abstract accumulator and reduce one-step public soundness to backend soundness, commitment binding, transcript binding, and row-level faithfulness of the typed implementation.
- **Implementation.** We map the construction to the current Symphony + WHIR codebase, emphasizing the public-proof boundary, the route split, and the concrete WHIR backend integration.
- **Evaluation.** We organize the performance story around the authoritative monolithic baseline, the explicit single-oracle accumulation route, and the integrated multi-oracle alternative.
- **Limitations.** We explain the current engineering boundary: the default public verifier is authoritative but still dominated by a large typed CP relation, and the next performance direction lies in more compressed SYMBT3/native multi-oracle structures.

2 Technical Overview

2.1 High-Level Architecture

At a high level, the repository implements a common stack and then exposes multiple public routes through that stack. The underlying mathematical objects are source R1CS/GR1CS-style claims together with Ajtai commitment openings. An individual WHIR proof certifies such a source claim. For accumulation, however, the system first views each statement/proof pair as a public WHIR proof-validity claim. Symphony’s high-arity folding machinery is applied to those claims, the resulting folding step is certified by a commit-and-prove relation, and that CP relation is proved together with the folded-output boundary consumed by the verifier.

A useful stack-level picture is shown in Figure 1.




Figure 1: High-level architecture of WHIR accumulation through lattice-based folding.

In the current implementation, WHIR is the main backend for this stack. It provides typed CP proofs and typed folded-output proofs over Poseidon2BabyBear public digests, while preserving a public-only verifier boundary for the authoritative public route.

2.2 What is Being Accumulated?

What is accumulated is not a vector of WHIR proof bytes. The accumulated object is a family of *WHIR proof-validity claims*. Each such claim includes the public inputs, the relation metadata, the public Fiat–Shamir commitments, the public roots and digests, the challenge binding, the commitment-opening algebra, and the folded-output semantics that the verifier must ultimately trust.

Thus, throughout the report, a claim means the public WHIR verifier obligation itself, not just a serialized proof object. Folding is sound only to the extent that it preserves that obligation and its public bindings.

2.3 Public-Verifier Boundary

Definition 2.1 (Authoritative public verifier). A verifier route is authoritative if its decision depends only on public inputs, public proof objects, public transcript commitments, public digests, and public backend proofs. It must not read witnesses, prover messages, Fiat–Shamir openings, fold inputs, folded witnesses, or debug-only replay material. Moreover, it must fail closed if the required typed backend verification is not available.

The authoritative public verifier is the default version-2 public proof route: `ProofBundleV2/PublicProofBundle` and `SymphonyProofV2/PublicSymphonyProof`, exposed through `prove_public` and `verify_public` with `prove_v2/verify_v2` as compatibility aliases. For WHIR+WHIR, this route is authoritative: the CP backend proves the typed CP relation, the output backend proves transcript binding for the folded output, the public digest scheme is `Poseidon2BabyBear`, and verification succeeds using public data only.

Concretely, the public verifier consumes caller-supplied public inputs and relation metadata, public Fiat–Shamir commitments, the public roots and digests `fs_root`, `fold_root`, `challenge_digest`, and `transcript_seed_digest`, the public folded-output instance, and the WHIR CP/output proofs. It must not consume Fiat–Shamir openings, Fiat–Shamir messages, fold inputs, original witnesses, folded witnesses, folding-proof internals, or CP witness bundles. Public verification is also fail-closed: if backend authority or typed verification is missing, the route rejects rather than silently falling back to witness-side checks.

This monolithic WHIR public verifier is therefore the correct baseline implementation for the report. It is the first unoptimized implementation stage in this codebase that establishes a real public-only boundary; every accumulation route studied later must preserve the same semantic commitments at that boundary.

2.4 Implementation Routes

The implementation routes relevant to this report are summarized in Table 1. They should be read as three implementation-level realizations of the same overall topic, each with a different role in the report:

- the authoritative public verifier provides the baseline public boundary;
- the single-oracle route provides the current explicit full-workload accumulation route;
- the multi-oracle route provides the integrated one-proof accumulation alternative.

This is a taxonomy, not a logical implication chain. The single-oracle and multi-oracle routes are explicit routes that are compared against the baseline public verifier; neither is a corollary of, nor a silent replacement for, the default public route.

Route	Conceptual role	Default?	Status
Product public verifier	Baseline public-only WHIR+WHIR verification boundary	Yes	Correct monolithic baseline
Single-oracle (repo K6a)	Explicit full-workload NonZK accumulation route	No	Implemented
Multi-oracle (repo N8)	Integrated same-shape NonZK accumulation route using one WHIR proof	No	Implemented, opt-in, not production-reviewed

Table 1: Verification and accumulation routes used in this report.

The authoritative default exposes `prove_public` and `verify_public`; `prove_v2/verify_v2` are compatibility aliases. The single-oracle route and the multi-oracle route are both explicit opt-in accumulation routes. The single-oracle route is the current main full-workload accumulation route; the multi-oracle route is the implemented integrated alternative with explicit `ACC.P/ACC.V/ACC.D` entry points. Their technical distinction is developed in Section 8.2. Neither silently replaces the default public verifier.

2.5 Performance Direction

The reason the story does not end at `verify_public` is performance. The default WHIR public verifier is authoritative and public-only, but it is still built around a monolithic typed CP relation whose row count grows materially with the folded workload. In other words, the baseline verifier is correct as a first unoptimized, monolithic implementation of the public boundary, but it is not yet the final cost profile one would want from an accumulated public route.

This motivates the broader SYMBT3/native multi-oracle direction. The report does not need to center every intermediate experiment in that line, but it does need to explain the performance motivation shared by the three routes above. The baseline public verifier establishes the current authoritative boundary; the single-oracle route is the current explicit accumulation realization studied against that baseline; and the multi-oracle route is the integrated one-WHIR alternative that explores a more compressed proof shape. The broader native/multi-oracle work should therefore be read mainly as the performance program that supports this line of refinement.

3 Preliminaries

We recall only the notation, interfaces, and soundness properties used later. The report treats WHIR, folding, commit-and-prove, and commitments as components with specified public behavior; their internal constructions are not reproved.

Notation. For $n \in \mathbb{N}$, write $[n] = \{1, \dots, n\}$. The security parameter is λ , and $\varepsilon(\lambda)$ denotes a negligible or explicitly bounded failure probability. Public inputs appear before a semicolon and private witness material after it:

$$R(x; w) = 1.$$

When no witness is shown, all inputs are public. We use $\text{Digest}(\cdot)$ for a collision-resistant or random-oracle-derived binding value, and H_{dom} for a domain-separated random-oracle call.

3.1 Interactive Oracle Proofs and SNARKs

An indexed relation is a family

$$R = \{R_{\text{idx}}\}_{\text{idx}},$$

where idx is a public index, x is a public instance, and w is a witness. We write

$$R_{\text{idx}}(x; w) = 1$$

when w is valid for x . Here, indices include proof-system parameters, relation identifiers, route metadata, and backend parameters.

A succinct argument system for R consists of a prover and verifier,

$$\Pi.\text{Prove}(\text{idx}, x, w) \rightarrow \pi, \quad \Pi.\text{Verify}(\text{idx}, x, \pi) \rightarrow \{0, 1\}.$$

Completeness means honest proofs verify. The knowledge-style soundness form used later says that an accepting proof implies, except with the stated error, the existence of a witness satisfying the indexed relation.

We use IOP/SNARK terminology only at this interface level. Oracle messages, Fiat–Shamir challenges, and openings matter in the report only insofar as they are committed public data that the verifier statement binds.

3.2 WHIR

WHIR [ACFY24b] is used as a transparent, hash-based backend for encoded relations:

$$\text{WHIR}.\text{Prove}(\text{stmt}, w) \rightarrow \Pi, \quad \text{WHIR}.\text{Verify}(\text{stmt}, \Pi) \rightarrow \{0, 1\}.$$

The public statement stmt contains the WHIR parameters, relation identifier, public inputs, public transcript material, and any folded-output boundary that the verifier is asked to trust.

We associate to WHIR a public verification relation

$$\mathcal{R}_{\text{WHIRVerify}}(\text{stmt}, \pi) = 1,$$

which holds exactly when the WHIR verifier accepts proof object π for stmt . The report uses WHIR in two roles:

1. input WHIR proofs π_i , which define the WHIR proof-validity claims being accumulated;
2. backend WHIR proofs Π_{cp} , which prove the encoded commit-and-prove relation certifying the folding transition.

Soundness is used as a black-box implication: if

$$\text{WHIR}.\text{Verify}(\text{stmt}, \Pi) = 1,$$

then, except with the WHIR soundness error, the relation encoded by stmt is satisfiable.

3.3 Accumulation and Folding

An accumulation scheme maintains a short public accumulator state. One step absorbs a batch of claims into a new state and proves that the update is valid. For WHIR proof-validity claims, the semantic folding transition is written

$$\mathcal{R}_{\text{Fold}}(\text{Acc}_{\text{old}}, \{\text{Claim}_i\}_{i=1}^k, \text{Acc}_{\text{new}}; w_{\text{fold}}) = 1.$$

The witness w_{fold} contains prover-side folding and opening material. The public verifier does not read it directly; it verifies a proof of the corresponding relation.

The relation enforces the supported claim shape, the prescribed linear or lattice-based combination of old state and claim data, and preservation of the public boundary: inputs, relation identifiers, transcript roots, challenge digests, and folded-output data all remain bound to the same transition. The Symphony-style layer is used only through this high-arity folding interface.

3.4 Commit-and-Prove SNARKs

A commit-and-prove relation checks ordinary algebraic constraints together with consistency against committed data. Here it bridges the semantic folding relation and the WHIR backend.

The public CP instance is written

$$I_{\text{cp}} = (\text{Acc}_{\text{old}}, \{\text{Claim}_i\}_{i=1}^k, \text{Acc}_{\text{new}}, \text{params}, \text{tr}_{\text{pub}}),$$

where tr_{pub} contains public transcript-binding material: relation identifiers, roots, digests, challenge digests, backend parameters, and folded-output boundary data. The witness is

$$w_{\text{cp}} = (w_{\text{fold}}, \text{openings}, \text{aux}).$$

The CP relation is denoted

$$\mathcal{R}_{\text{CPFold}}(I_{\text{cp}}; w_{\text{cp}}) = 1.$$

The role of $\mathcal{R}_{\text{CPFold}}$ is to certify the folding transition without exposing folding witnesses or openings. It must not weaken the folding statement: whenever

$$\mathcal{R}_{\text{CPFold}}(I_{\text{cp}}; w_{\text{cp}}) = 1,$$

the underlying folding relation must hold for the same public claims, the same transcript-binding material, and the same folded-output boundary.

In the implementation, this CP relation is encoded as a WHIR backend statement and proved using WHIR; there is no separate CP proof followed by a WHIR proof.

3.5 Lattice Commitments

The folding layer uses lattice/Ajtai-style commitments to bind folding data and witness material. We use only the interface

$$c \leftarrow \text{Commit}(m; r),$$

where m is the committed message and r is opening randomness or opening material.

The property used later is binding: except with error $\varepsilon_{\text{bind}}$, a commitment cannot be opened to two inconsistent messages. Binding ties folding witnesses, claim encodings, accumulator state, and CP witness material to the same committed data.

There are several commitment layers in the construction. Lattice commitments bind folding data. The commit-and-prove layer binds witness material to the public CP instance. WHIR uses hash/Merkle-style oracle commitments internally when proving the encoded backend relation. These layers serve different roles and must be tied to the same public statement by explicit transcript and digest binding.

4 WHIR Accumulation Through Lattice-Based Folding

4.1 Interface

Algorithm 1 ACCUMULATEWHIR

Require: Old accumulator Acc_{old}

Require: Public WHIR proof-validity inputs represented by an accumulation batch

Ensure: New accumulator Acc_{new} and accumulation proof Π_{acc}

- 1: Convert the batch inputs into WHIR proof-validity claims.
 - 2: Fold the resulting claims using the lattice-based folding layer.
 - 3: Form the CP instance I_{cp} and witness w_{cp} certifying the folding transition.
 - 4: Encode $\mathcal{R}_{\text{CPFold}}(I_{\text{cp}}; w_{\text{cp}}) = 1$ as a WHIR backend statement stmt_{cp} .
 - 5: Produce $\Pi_{\text{cp}} \leftarrow \text{WHIR.Prove}(\text{stmt}_{\text{cp}}, w_{\text{cp}})$.
 - 6: **return** $(\text{Acc}_{\text{new}}, \Pi_{\text{acc}})$, where Π_{acc} contains Π_{cp} and the public CP instance.
-

4.2 From WHIR Proofs to Verification Claims

We now fix notation for the claim object introduced above. A WHIR prover produces a proof object π_i for a public verifier statement stmt_i ; accumulation uses the corresponding verifier predicate:

$$(\text{stmt}_i, \pi_i) \rightsquigarrow \mathcal{R}_{\text{WHIRVerify}}(\text{stmt}_i, \pi_i) = 1.$$

Definition 4.1 (WHIR proof-validity claim). A WHIR proof-validity claim is a pair (stmt, π) together with the assertion

$$\mathcal{R}_{\text{WHIRVerify}}(\text{stmt}, \pi) = 1,$$

where stmt is the full public WHIR verifier statement and π is the corresponding WHIR proof object.

The statement stmt_i is the full public statement at the WHIR verifier boundary, not merely the original source relation. In particular, it contains the declared WHIR parameters, relation metadata, public inputs, public transcript commitments, roots and digests, challenge-binding data, and the folded-output boundary that the verifier is asked to trust.

The accumulator therefore acts on $\mathcal{R}_{\text{WHIRVerify}}(\text{stmt}_i, \pi_i) = 1$ together with these public bindings. Later soundness statements rely on the folded claim remaining the same verifier obligation as the ordinary WHIR relation $\mathcal{R}_{\text{WHIRVerify}}$.

Equivalently, for a batch of WHIR statements and proofs we write

$$\text{Claim}_i := (\text{stmt}_i, \pi_i) \quad \text{with predicate} \quad \mathcal{R}_{\text{WHIRVerify}}(\text{stmt}_i, \pi_i) = 1, \quad i \in [k].$$

The accumulation procedure acts on the family $\{\text{Claim}_i\}_{i=1}^k$. Later sections specify how these claims are folded, how the folding transition is certified by a commit-and-prove relation, and how that relation is encoded into the WHIR backend. The soundness of the overall construction depends on this encoding proving the same public relation.

4.3 Lattice-Based Folding of Claims

After converting WHIR statement/proof pairs into proof-validity claims, the accumulator must combine a batch of such claims into one transition. Let $\text{Claim}_1, \dots, \text{Claim}_k$ be the WHIR proof-validity claims defined in Section 4.2. The role of the folding layer is to absorb the family $\{\text{Claim}_i\}_{i=1}^k$ into a transition from an old accumulator Acc_{old} to a new accumulator Acc_{new} .

At this level, the folding layer should be understood as an accumulator update relation. It does not merely hash together the input proof objects, nor does it verify each WHIR proof independently. Instead, it constructs a folded public instance whose validity is intended to imply that the input proof-validity claims have been correctly incorporated into the accumulator state.

We write the abstract folding relation as

$$\mathcal{R}_{\text{Fold}}(\text{Acc}_{\text{old}}, \{\text{Claim}_i\}_{i=1}^k, \text{Acc}_{\text{new}}; w_{\text{fold}}) = 1,$$

where w_{fold} denotes the private material needed to witness the folding transition. In the implementation this material may include commitment openings, folding witnesses, and other prover-side data needed to justify the new folded state. The public verifier will not read this material directly; it will only check a proof that the corresponding relation is satisfied.

At this point $\mathcal{R}_{\text{Fold}}$ is an abstract relation: it specifies the semantic accumulator transition, not yet the concrete arithmetized relation proved by WHIR.

The folding transition is required to bind the following public data:

1. the old accumulator Acc_{old} ;
2. the full batch of WHIR proof-validity claims $\{\text{Claim}_i\}_{i=1}^k$;
3. the declared relation and backend parameters;
4. the Fiat–Shamir transcript material used to derive the folding challenges;
5. the folded-output boundary that becomes part of Acc_{new} .

Concretely, the folding challenge vector should be derived from public binding material, rather than chosen independently by the prover. Schematically, one may view it as

$$\beta := H_{\text{fold}}(\text{Acc}_{\text{old}}, \text{Digest}(\text{Claim}_1, \dots, \text{Claim}_k), \text{params}),$$

where H_{fold} denotes the domain-separated random-oracle call used for folding. The exact transcript format is implementation-dependent, but the soundness requirement is not: the same public material that determines the claims must also determine the folding challenges.

The relation $\mathcal{R}_{\text{Fold}}$ enforces three kinds of consistency.

First, it enforces *shape consistency*: the claims being folded must belong to the same supported accumulation shape. In the implemented NonZK routes this means, in particular, that the accumulation batch is nonempty and that the claims have compatible relation and backend metadata.

Second, it enforces *algebraic folding consistency*: the new folded state must be the prescribed linear or lattice-based combination of the previous state and the input claim data under the challenge vector β . This is the part of the construction inherited from the Symphony-style folding layer.

Third, it enforces *public-boundary consistency*: the accumulator update must preserve the same public statement data that the verifier is asked to trust. The folded accumulator must therefore be bound to the same public inputs, relation identifiers, transcript roots, challenge digests, and folded-output boundary that appear in the input WHIR proof-validity claims.

The output of this stage is not yet a public proof. It is a folded transition claim:

$$\mathcal{R}_{\text{Fold}}(\text{Acc}_{\text{old}}, \{\text{Claim}_i\}_{i=1}^k, \text{Acc}_{\text{new}}; w_{\text{fold}}) = 1.$$

The next section explains how this semantic transition relation is turned into a commit-and-prove relation whose public instance is suitable for backend verification.

4.4 CP-SNARK for the Folding Transition

The folding relation defined above describes the semantic accumulator update, but it is not yet a public proof object. The verifier should not receive the folding witness w_{fold} , the commitment openings, or any replay material used by the prover. Instead, the folding transition is packaged as a commit-and-prove relation.

We write the public CP instance as

$$I_{\text{cp}} := (\text{Acc}_{\text{old}}, \{\text{Claim}_i\}_{i=1}^k, \text{Acc}_{\text{new}}, \text{params}, \text{tr}_{\text{pub}}),$$

where tr_{pub} denotes the public transcript-binding material: relation identifiers, public roots and digests, challenge digests, backend parameters, and the folded-output boundary. The corresponding witness is

$$w_{\text{cp}} := (w_{\text{fold}}, \text{openings}, \text{aux}),$$

where openings denotes the commitment-opening material and aux denotes any auxiliary prover-side data needed to check the folding transition.

The CP relation is then written as

$$\mathcal{R}_{\text{CPFold}}(I_{\text{cp}}; w_{\text{cp}}) = 1.$$

Its purpose is to certify that the transition from Acc_{old} to Acc_{new} is a valid folding transition for the declared batch of WHIR proof-validity claims. In particular, $\mathcal{R}_{\text{CPFold}}$ checks that:

1. the public instance contains the same claims, parameters, transcript digests, and folded-output boundary that the verifier will consume;
2. the folding challenges are derived from the declared public transcript material using the correct domain separation;
3. the private openings and folding witness are consistent with the committed claim data;
4. the algebraic folding equations of $\mathcal{R}_{\text{Fold}}$ hold;
5. the resulting folded state is exactly the declared accumulator Acc_{new} .

Thus $\mathcal{R}_{\text{CPFold}}$ is the public-verifier version of the folding transition. It hides the prover-side witness material, but it must not weaken the statement being proved. Soundness requires that whenever

$$\mathcal{R}_{\text{CPFold}}(I_{\text{cp}}; w_{\text{cp}}) = 1,$$

the underlying folding relation

$$\mathcal{R}_{\text{Fold}}(\text{Acc}_{\text{old}}, \{\text{Claim}_i\}_{i=1}^k, \text{Acc}_{\text{new}}; w_{\text{fold}}) = 1$$

also holds for the same public claims, the same transcript-binding material, and the same folded-output boundary.

This is the point at which commit-and-prove is essential. The CP layer lets the prover commit to, and prove consistency of, the folding witness and opening material without exposing those objects to the public verifier. The verifier will later check only a succinct backend proof of $\mathcal{R}_{\text{CPFold}}$, together with the public instance I_{cp} .

4.5 WHIR Backend for the CP Relation

The CP relation from the previous section is still only a relation. To obtain a succinct public proof, the construction uses WHIR as the backend proof system for this relation. This is the second role played by WHIR in the construction: the input objects are WHIR proof-validity claims, while the backend proof is a new WHIR proof attesting to the CP relation that certifies the folding transition.

We write the WHIR encoding of the CP instance as

$$\text{stmt}_{\text{cp}} := \text{Enc}_{\text{WHIR}}(\mathcal{R}_{\text{CPFold}}, I_{\text{cp}}, \text{pp}_{\text{WHIR}}),$$

where pp_{WHIR} denotes the public WHIR parameters used by the backend. The backend prover then produces a WHIR proof

$$\Pi_{\text{cp}} \leftarrow \text{WHIR.Prove}(\text{stmt}_{\text{cp}}, w_{\text{cp}}).$$

The public verifier checks this proof by running

$$\text{WHIR.Verify}(\text{stmt}_{\text{cp}}, \Pi_{\text{cp}}) = 1.$$

The object Π_{cp} is the backend proof carried by the accumulator. It should not be confused with the input WHIR proofs $\pi_1, \dots, \pi_k$. Those input proofs define the proof-validity claims being accumulated; the backend proof Π_{cp} certifies that the folding transition over those claims was performed correctly.

The WHIR backend must bind the same public data as the CP instance. In particular, stmt_{cp} must include, or be bound to, the old accumulator, the new accumulator, the batch of WHIR proof-validity claims, the declared relation parameters, the transcript-binding material, and the folded-output boundary. If any of these values are omitted from the backend statement, then the backend proof may be sound for a weaker relation than the one required by the accumulator.

The intended implication of the backend layer is

$$\text{WHIR.Verify}(\text{stmt}_{\text{cp}}, \Pi_{\text{cp}}) = 1 \implies \exists w_{\text{cp}} : \mathcal{R}_{\text{CPFold}}(I_{\text{cp}}; w_{\text{cp}}) = 1,$$

up to the soundness error of the WHIR backend. This implication is the point at which backend soundness enters the overall argument. The next soundness step is then supplied by the CP relation itself:

$$\mathcal{R}_{\text{CPFold}}(I_{\text{cp}}; w_{\text{cp}}) = 1 \implies \mathcal{R}_{\text{Fold}}(\text{Acc}_{\text{old}}, \{\text{Claim}_i\}_{i=1}^k, \text{Acc}_{\text{new}}; w_{\text{fold}}) = 1.$$

This separation is important. WHIR soundness alone only says that the encoded backend statement is satisfied. It does not by itself say that the accumulator transition is sound. The composition is sound only if the WHIR statement stmt_{cp} faithfully encodes the CP relation $\mathcal{R}_{\text{CPFold}}$ for the same public instance I_{cp} . This faithfulness condition is stated explicitly in Section 5.5.

Thus the backend layer has a precise role: it turns the CP folding claim into a succinct public proof. The semantic content of the accumulator remains in $\mathcal{R}_{\text{Fold}}$ and $\mathcal{R}_{\text{CPFold}}$; WHIR supplies the succinct verification mechanism for that content.

4.6 Verification Algorithm

We now describe the public verification algorithm for an accumulated WHIR proof. The verifier receives the old accumulator, the new accumulator, the batch of public WHIR proof-validity claims, and the backend WHIR proof of the CP relation. It does not receive the folding witness, the commitment openings, the original witnesses, or any prover-side replay material.

The public accumulation proof has the form

$$\Pi_{\text{acc}} := (I_{\text{cp}}, \text{stmt}_{\text{cp}}, \Pi_{\text{cp}}),$$

where I_{cp} is the public CP instance, stmt_{cp} is the WHIR encoding of that instance, and Π_{cp} is the WHIR backend proof.

The verifier performs the following checks.

1. **Shape and boundary checks.** The verifier checks that I_{cp} contains the declared old accumulator Acc_{old} , the declared new accumulator Acc_{new} , the declared batch $\{\text{Claim}_i\}_{i=1}^k$, and the expected public parameters.
2. **Transcript-binding checks.** The verifier recomputes the public transcript-binding material from the declared claims and parameters. In particular, the public roots, digests, challenge digest, and folded-output boundary must agree with the values included in I_{cp} .
3. **Backend statement check.** The verifier checks that the WHIR backend statement is the prescribed encoding of the CP relation and public instance:

$$\text{stmt}_{\text{cp}} = \text{Enc}_{\text{WHIR}}(\mathcal{R}_{\text{CPFold}}, I_{\text{cp}}, \text{pp}_{\text{WHIR}}).$$

4. **WHIR backend verification.** Finally, the verifier runs

$$\text{WHIR.Verify}(\text{stmt}_{\text{cp}}, \Pi_{\text{cp}}).$$

The accumulation proof is accepted if and only if this check accepts.

Equivalently, the public verifier can be written as

$$\text{AccVerify}(\text{Acc}_{\text{old}}, \{\text{Claim}_i\}_{i=1}^k, \text{Acc}_{\text{new}}, \Pi_{\text{acc}}) = 1$$

if and only if all public boundary checks succeed and

$$\text{WHIR.Verify}(\text{stmt}_{\text{cp}}, \Pi_{\text{cp}}) = 1.$$

The fail-closed condition is important. If the backend statement cannot be recomputed from the declared public instance, if the public transcript material does not match, or if the typed WHIR backend verification is unavailable, the verifier rejects. In particular, the verifier must not repair the transcript using witness-side data, replay prover messages, or fall back to a non-authoritative helper.

This completes the construction. The next section isolates the relation-level definitions needed for the completeness and soundness arguments. The most important of these is the faithfulness condition: the WHIR backend statement must encode exactly the CP relation required by the folding transition, not a weaker or differently-bound relation.

5 Formal Relation Definitions

Section 4 described the construction operationally. We now collect the relations that are used in the completeness and soundness arguments. The purpose of this section is to make explicit what is verified at each layer, and in particular to separate four different notions:

1. validity of an input WHIR proof;
2. validity of the folding transition that accumulates such claims;
3. validity of the commit-and-prove relation certifying the transition;
4. validity of the WHIR backend proof for the CP relation.

5.1 WHIR Verification Relation

A WHIR verifier statement is written

$$\text{stmt} = (\text{pp}_{\text{WHIR}}, \text{rel_id}, x_{\text{pub}}, \text{tr}_{\text{pub}}, \text{boundary}),$$

where pp_{WHIR} are the public WHIR parameters, rel_id identifies the proved relation, x_{pub} denotes the public inputs, tr_{pub} contains the public transcript roots and digests, and boundary denotes the folded-output boundary that the verifier is asked to trust.

The WHIR verification relation is the public predicate

$$\mathcal{R}_{\text{WHIRVerify}}(\text{stmt}, \pi) = 1$$

that holds if and only if the WHIR verifier accepts π for stmt . Semantically, this relation checks that:

1. the declared WHIR parameters and relation identifier are well formed;
2. the public transcript commitments, roots, and digests are consistent with the proof object;
3. the Fiat–Shamir challenges are derived from the declared public transcript using the correct domain separation;
4. the opened evaluations and query answers are consistent with the committed oracles;
5. the encoded source relation and folded-output boundary checks required by the WHIR statement are satisfied.

A WHIR proof-validity claim is valid precisely when

$$\mathcal{R}_{\text{WHIRVerify}}(\text{stmt}_i, \pi_i) = 1.$$

5.2 Folding Relation

Let

$$\text{Claim}_i := (\text{stmt}_i, \pi_i) \quad \text{with predicate} \quad \mathcal{R}_{\text{WHIRVerify}}(\text{stmt}_i, \pi_i) = 1.$$

The folding relation is written

$$\mathcal{R}_{\text{Fold}}(\text{Acc}_{\text{old}}, \{\text{Claim}_i\}_{i=1}^k, \text{Acc}_{\text{new}}; w_{\text{fold}}) = 1.$$

This relation is the semantic accumulator-transition relation. It is not merely a check that the proof objects have been hashed into the accumulator. It asserts that the batch of WHIR proof-validity claims has been correctly absorbed into the transition from Acc_{old} to Acc_{new} .

Concretely, $\mathcal{R}_{\text{Fold}}$ checks the following conditions.

1. **Claim-validity representation.** The folded relation represents the validity of the input WHIR proof-validity claims. Semantically, an accepting accumulator transition should imply

$$\forall i \in [k], \quad \mathcal{R}_{\text{WHIRVerify}}(\text{stmt}_i, \pi_i) = 1.$$

The implemented single-oracle and multi-oracle routes need not realize this by embedding a generic WHIR verifier for arbitrary proof objects. Instead, they represent the relevant validity obligations through typed SYMBT3/K6a/N8 semantic rows, transition rows, and binding rows. The soundness requirement is that these rows faithfully encode the WHIR-backed validity obligations being accumulated.

2. **Shape consistency.** The claims belong to the same supported accumulation shape: relation metadata, backend parameters, field/ring parameters, arity, and folded-output format are compatible.
3. **Challenge binding.** The folding challenge vector β is derived from the declared public transcript material, old accumulator, claim digest, and backend parameters using the prescribed domain separation.
4. **Commitment and opening consistency.** The private opening material in w_{fold} is consistent with the public commitments that appear in the accumulator transition.
5. **Algebraic folding consistency.** The new folded state is the prescribed linear or lattice-based combination of the old state and the input claim data under β .
6. **Output consistency.** The computed folded state is exactly the declared accumulator Acc_{new} , including the folded-output boundary that the public verifier later consumes.

5.3 Commit-and-Prove Folding Relation

The public CP instance is

$$I_{\text{cp}} := \left(\text{Acc}_{\text{old}}, \{\text{Claim}_i\}_{i=1}^k, \text{Acc}_{\text{new}}, \text{params}, \text{tr}_{\text{pub}} \right),$$

and the CP witness is

$$w_{\text{cp}} := (w_{\text{fold}}, \text{openings}, \text{aux}).$$

The commit-and-prove folding relation is written

$$\mathcal{R}_{\text{CPFold}}(I_{\text{cp}}; w_{\text{cp}}) = 1.$$

It holds when the witness material opens the committed folding data correctly and when the semantic folding relation holds:

$$\mathcal{R}_{\text{Fold}}\left(\text{Acc}_{\text{old}}, \{\text{Claim}_i\}_{i=1}^k, \text{Acc}_{\text{new}}; w_{\text{fold}}\right) = 1,$$

with the same public claims, transcript material, parameters, and folded-output boundary as those appearing in I_{cp} .

Thus $\mathcal{R}_{\text{CPFold}}$ is not a weaker replacement for $\mathcal{R}_{\text{Fold}}$. It is the committed, public-verifier form of the same accumulator-transition statement.

5.4 WHIR Backend Relation

The WHIR backend encodes the CP relation into a WHIR statement

$$\text{stmt}_{\text{cp}} := \text{Enc}_{\text{WHIR}}(\mathcal{R}_{\text{CPFold}}, I_{\text{cp}}, \text{pp}_{\text{WHIR}}).$$

The backend proof is Π_{cp} , and the backend verification predicate is

$$\mathcal{R}_{\text{backend}}(\text{stmt}_{\text{cp}}, \Pi_{\text{cp}}) = 1 \iff \text{WHIR.Verify}(\text{stmt}_{\text{cp}}, \Pi_{\text{cp}}) = 1.$$

Soundness of the backend gives the implication

$$\mathcal{R}_{\text{backend}}(\text{stmt}_{\text{cp}}, \Pi_{\text{cp}}) = 1 \implies \exists w_{\text{cp}} : \mathcal{R}_{\text{CPFold}}(I_{\text{cp}}; w_{\text{cp}}) = 1,$$

except with the soundness error of the WHIR backend and by the soundness direction of the faithfulness obligation stated in Section 5.5.

5.5 Faithfulness of the Encoded Relation

The WHIR backend is sound for the relation that it actually encodes. Therefore, the composition is sound only if the encoded backend statement is faithful to the CP relation required by the accumulator.

We say that the WHIR encoding is faithful if, for every public CP instance I_{cp} and witness w_{cp} ,

$$\mathcal{R}_{\text{CPFold}}(I_{\text{cp}}; w_{\text{cp}}) = 1 \iff \mathcal{R}_{\text{enc}}(\text{Enc}_{\text{WHIR}}(I_{\text{cp}}), \text{Enc}_{\text{WHIR}}(w_{\text{cp}})) = 1,$$

where $\mathcal{R}_{\text{enc}}$ is the relation represented by the WHIR backend statement.

Equivalently, the WHIR backend must prove exactly the CP folding relation for the same public instance: the same old accumulator, the same new accumulator, the same WHIR proof-validity claims, the same transcript-binding material, and the same folded-output boundary.

If the backend statement omits one of these public values, proves the folding relation under different transcript material, or accepts a weaker relation, then WHIR may remain sound for its own encoded statement while the accumulator is sound for the wrong language. Faithfulness is the condition that rules out this gap.

5.6 Accumulator Validity and Concrete Route Faithfulness

The relation definitions above are abstract: they use a family of WHIR proof-validity claims

$$\{\text{Claim}_i\}_{i=1}^k$$

as the input to the folding relation. The implementation, however, does not expose an unrestricted vector of arbitrary WHIR statement/proof pairs. It exposes typed accumulation batches for the single-oracle and multi-oracle routes. We therefore make explicit the bridge between the concrete route and the abstract relation.

We write

$$\text{ValidAcc}(\text{Acc})$$

for the accumulator invariant. Informally, this means that Acc is a well-formed accumulator state for the declared route, relation shape, backend parameters, transcript domain, and folded-output boundary. The exact well-formedness conditions are route-specific, but they must include the public metadata consumed by the verifier.

For example, in the K6a route, **ValidAcc** means that the public accumulator instance and SYMBT3 backend public statement contain well-formed old and new accumulator digests and coordinates, the K6a shape and profile digests, WHIR/BabyBear parameter digests, manifest and source layout digests, source assignment and Ajtai-opening root digests, message-oracle root digests, and the folded-output boundary, all bound to the NonZK-integrity K6a relation accepted by the route policy. For the N8 integrated route, **ValidAcc** additionally includes the native multi-oracle descriptor data: version and workload kind, the K6a relation identifier, tuple-leaf descriptor and layout digests, same-field/rate/folding-parameter flags, claim-plan digest, committed table digest, semantic-completion flags, and N8 transcript-binding digest. These examples are not new assumptions; they spell out what the route-specific well-formed public accumulator invariant must contain.

Let B denote a concrete typed accumulation batch. A route-specific interpretation map

$$\text{RouteMap}(B) = \{\text{Claim}_i\}_{i=1}^k$$

extracts the abstract WHIR-backed proof-validity obligations represented by that batch. This map is not required to expose a literal list of arbitrary WHIR proof objects. It only states the semantic claims represented by the typed rows of the route.

We write

$$\text{ConcreteRoute}(B, \text{Acc}_{\text{old}}, \text{Acc}_{\text{new}}; \omega) = 1$$

for the concrete row-level relation implemented by a route, where ω denotes the private route witness, including openings, auxiliary folding material, and route-specific prover data.

Definition 5.1 (Typed-batch faithfulness). A concrete route is faithful to the abstract folding relation if, for every typed batch B , old accumulator Acc_{old} , new accumulator Acc_{new} , and route witness ω ,

$$\text{ConcreteRoute}(B, \text{Acc}_{\text{old}}, \text{Acc}_{\text{new}}; \omega) = 1$$

implies that, for

$$\{\text{Claim}_i\}_{i=1}^k := \text{RouteMap}(B),$$

there exists a folding witness w_{fold} such that

$$\mathcal{R}_{\text{Fold}}(\text{Acc}_{\text{old}}, \{\text{Claim}_i\}_{i=1}^k, \text{Acc}_{\text{new}}; w_{\text{fold}}) = 1.$$

Typed-batch faithfulness is the formal version of the implementation claim that the single-oracle and multi-oracle routes represent WHIR-backed proof-validity obligations through typed SYMBT3/K6a/N8 rows. It is not automatic. It must be checked by comparing the route rows against the abstract relation ledger: claim-validity representation, shape consistency, challenge binding, commitment opening consistency, algebraic folding consistency, and output consistency.

For completeness, one also wants the converse direction for honest route generation: if the abstract WHIR-backed claims are valid and the honest route prover follows the construction, then the concrete route relation is satisfied. For soundness, the direction above is the critical one.

5.7 Row-Level Faithfulness Obligations

Definition 5.1 states typed-batch faithfulness as an abstract condition. For the implemented routes, the remaining mathematical task is to prove this condition from the concrete rows generated by the backend statement. This subsection isolates the row-level obligations that such a proof must discharge.

Let B be a typed accumulation batch and let

$$\text{RouteMap}(B) = \{\text{Claim}_i\}_{i=1}^k.$$

Let

$$\text{Rows}(B)$$

denote the concrete row system encoded into the backend WHIR statement. We decompose these rows conceptually as

$$\text{Rows}(B) = \text{Rows}_{\text{claim}} \cup \text{Rows}_{\text{shape}} \cup \text{Rows}_{\text{challenge}} \cup \text{Rows}_{\text{opening}} \cup \text{Rows}_{\text{fold}} \cup \text{Rows}_{\text{output}}.$$

This decomposition is logical rather than necessarily syntactic: a concrete row may contribute to more than one of these checks.

Lemma 5.2 (Sufficient conditions for typed-row faithfulness). *Suppose that for every typed batch B , old accumulator Acc_{old} , new accumulator Acc_{new} , and concrete route witness ω , satisfaction of the concrete row system implies the following six properties.*

1. **Claim representation.** *The claim rows identify a unique family of represented WHIR-backed validity obligations*

$$\{\text{Claim}_i\}_{i=1}^k := \text{RouteMap}(B),$$

and bind those obligations to the same public metadata, transcript roots, challenge digests, folded-output boundary, and batch digest used by the public verifier. This condition does not mean that the rows run a generic WHIR verifier for each input proof; it means that the typed row system represents the same public obligations that the abstract claims denote.

2. **Shape consistency.** *The shape rows enforce that all represented claims have compatible relation metadata, backend parameters, field or ring parameters, arity, and folded-output format.*
3. **Challenge binding.** *The challenge rows enforce that the folding challenge vector β is derived from the declared public transcript material, old accumulator, claim digest, and backend parameters using the prescribed domain separation.*
4. **Opening consistency.** *The opening rows enforce that every private opening used by the folding witness is consistent with the public commitments appearing in the typed batch and accumulator transition.*
5. **Algebraic folding consistency.** *The folding rows enforce that the new folded state is exactly the prescribed linear or lattice-based combination of the old state and the represented claim data under β .*
6. **Output consistency.** *The output rows enforce that the computed folded state is exactly the public accumulator Acc_{new} , including the folded-output boundary later consumed by the verifier.*

Then the concrete route is faithful to the abstract folding relation:

$$\begin{aligned} \text{ConcreteRoute}(B, \text{Acc}_{\text{old}}, \text{Acc}_{\text{new}}; \omega) &= 1 \\ \implies \exists w_{\text{fold}} : \mathcal{R}_{\text{Fold}}(\text{Acc}_{\text{old}}, \text{RouteMap}(B), \text{Acc}_{\text{new}}; w_{\text{fold}}) &= 1. \end{aligned}$$

Proof. Assume that the concrete route relation accepts. By the claim-representation rows, the typed batch determines the abstract family of represented WHIR-backed proof-validity claims. By the shape rows, these claims belong to a supported accumulation shape. By the challenge rows, the folding challenge vector used by the concrete route is the one prescribed by the public transcript. By the opening rows, the private witness material used by the route is consistent with the public commitments. By the folding rows, the new folded state is the prescribed algebraic combination of the old accumulator and represented claim data. Finally, by the output rows, this computed folded state is the declared public accumulator Acc_{new} . These are exactly the conditions listed in the abstract folding relation $\mathcal{R}_{\text{Fold}}$. Hence there exists a folding witness w_{fold} , extracted from the concrete route witness ω , such that the abstract folding relation holds. $\square$

The claim-representation gap. The first condition in Lemma 5.2 is the load-bearing one. The implemented K6a and N8 rows do not embed a generic WHIR verifier for arbitrary statement/proof pairs. Instead, they represent WHIR-backed validity obligations through typed semantic rows, transition rows, tuple-RLC rows, opening rows, and binding rows.

Thus the lemma should be read as reducing concrete soundness to the claim-representation condition, modulo the easier five structural conditions. Property (1) is the deepest part of the six: it is where the concrete typed rows must be shown to denote the same WHIR-backed validity claims as the abstract relation, rather than merely to carry similarly named digests.

Consequently, the statement

$$\text{RouteMap}(B) = \{\text{Claim}_i\}_{i=1}^k$$

should not be read as extracting a literal vector of arbitrary WHIR proofs from the batch. It is an interpretation map: it says which public verification obligations the typed row system is meant to represent. To prove concrete soundness for K6a or N8, one must show that the row system enforces those obligations with the same public metadata, transcript roots, challenge digests, and folded-output boundary as the ordinary WHIR verification relation.

This is plausible because the typed rows are designed around precisely those public objects, but it is not proved in this report. A complete row-level proof would have to show, for each row family, that satisfying the rows implies the corresponding component of the represented WHIR-backed verification obligation. In particular, it would have to rule out the following failure mode: the row system binds a digest of a proof object, or a digest of a semantic claim, without binding the full public verifier relation that makes the claim meaningful.

This lemma does not by itself prove that the current implementation satisfies the six row obligations. Rather, it reduces the remaining implementation-specific proof to a row-by-row audit. For K6a and N8, such an audit must show that the SYMBT3/K6a/N8 semantic rows, transition rows, tuple-RLC rows, opening rows, and binding rows jointly imply the six properties above. This is the largest remaining theoretical gap between the abstract soundness theorem and the concrete implementation.

6 Completeness

We now prove completeness of the abstract accumulator. Completeness is the honest-prover direction: if the input batch represents valid WHIR-backed proof-validity claims, if the folding prover constructs the accumulator update according to the prescribed relation, and if the WHIR backend is run honestly on the encoded CP relation, then the public accumulation verifier accepts.

The statement is conditional because the notation $\{\text{Claim}_i\}_{i=1}^k$ is a semantic abstraction. The implemented single-oracle and multi-oracle routes consume typed accumulation batches and route-specific public instances, not an unrestricted vector of arbitrary WHIR proofs. Completeness is therefore stated for any concrete route whose typed batch faithfully represents the abstract claims and whose prover follows the construction of Section 4.

6.1 Local Completeness of the WHIR Claims

Let

$$\text{Claim}_i := (\text{stmt}_i, \pi_i), \quad i \in [k],$$

be the WHIR proof-validity claims represented by the typed input batch. Local completeness of the input claims means that every represented WHIR proof verifies against its declared public statement:

$$\forall i \in [k], \quad \mathcal{R}_{\text{WHIRVerify}}(\text{stmt}_i, \pi_i) = 1.$$

This is the local validity condition for the objects being accumulated. In the concrete implementation, these claims may be represented through typed SYMBT3/K6a/N8 public structures rather than through a literal $\text{Vec} < (\text{stmt}, \pi) >$. The semantic requirement is that the typed batch denotes valid WHIR-backed verification obligations.

6.2 Completeness of Folding

Assume that the old accumulator Acc_{old} is valid for the declared accumulation shape and that the represented input claims are valid and shape-compatible. The honest folding prover derives the folding challenges from the prescribed public transcript material and computes the new accumulator Acc_{new} together with folding witness material w_{fold} .

By completeness of the folding procedure, the resulting transition satisfies

$$\mathcal{R}_{\text{Fold}}(\text{Acc}_{\text{old}}, \{\text{Claim}_i\}_{i=1}^k, \text{Acc}_{\text{new}}; w_{\text{fold}}) = 1.$$

This says that the represented input claims have been correctly absorbed into the new accumulator state, with the declared shape, transcript binding, folding challenges, and folded-output boundary.

6.3 Completeness of the CP Layer

The honest prover next forms the public CP instance

$$I_{\text{cp}} := (\text{Acc}_{\text{old}}, \{\text{Claim}_i\}_{i=1}^k, \text{Acc}_{\text{new}}, \text{params}, \text{tr}_{\text{pub}}),$$

and the CP witness

$$w_{\text{cp}} := (w_{\text{fold}}, \text{openings}, \text{aux}).$$

If the folding transition is valid and the commitment-opening material is correct, then the commit-and-prove folding relation is satisfied:

$$\mathcal{R}_{\text{CPFold}}(I_{\text{cp}}; w_{\text{cp}}) = 1.$$

Thus the CP layer preserves the honest folding transition while keeping the folding witness, openings, and auxiliary prover-side material outside the public verifier boundary.

6.4 Completeness of the WHIR Backend

The CP instance is encoded as a WHIR backend statement

$$\text{stmt}_{\text{cp}} := \text{Enc}_{\text{WHIR}}(\mathcal{R}_{\text{CPFold}}, I_{\text{cp}}, \text{pp}_{\text{WHIR}}).$$

The honest backend prover computes

$$\Pi_{\text{cp}} \leftarrow \text{WHIR.Prove}(\text{stmt}_{\text{cp}}, w_{\text{cp}}).$$

By completeness of the WHIR backend, if

$$\mathcal{R}_{\text{CPFold}}(I_{\text{cp}}; w_{\text{cp}}) = 1,$$

then the public backend verifier accepts:

$$\text{WHIR.Verify}(\text{stmt}_{\text{cp}}, \Pi_{\text{cp}}) = 1.$$

6.5 Full Completeness Theorem

Theorem 6.1 (Conditional completeness of the abstract accumulator). *Assume the following:*

1. the typed input batch denotes a family of valid WHIR-backed proof-validity claims $\{\text{Claim}_i\}_{i=1}^k$;
2. the old accumulator Acc_{old} is valid for the declared accumulation shape;
3. the honest folding prover computes Acc_{new} and w_{fold} according to the prescribed folding relation;
4. the CP instance I_{cp} and witness w_{cp} are formed from the same public claims, transcript material, parameters, and folded-output boundary;
5. the concrete route is honest-complete in the converse direction of Definition 5.1: an honestly generated valid abstract folding transition satisfies the concrete route relation;
6. the WHIR backend is complete for the encoded CP relation.

Then the public accumulation verifier accepts:

$$\text{AccVerify}(\text{Acc}_{\text{old}}, \{\text{Claim}_i\}_{i=1}^k, \text{Acc}_{\text{new}}, \Pi_{\text{acc}}) = 1.$$

Proof. By assumption, the typed input batch denotes valid WHIR-backed proof-validity claims. Semantically, this means that for every represented claim,

$$\mathcal{R}_{\text{WHIRVerify}}(\text{stmt}_i, \pi_i) = 1.$$

Because the old accumulator is valid and the claims have the declared supported shape, completeness of the folding procedure gives a witness w_{fold} such that

$$\mathcal{R}_{\text{Fold}}(\text{Acc}_{\text{old}}, \{\text{Claim}_i\}_{i=1}^k, \text{Acc}_{\text{new}}; w_{\text{fold}}) = 1.$$

The honest prover forms the CP instance I_{cp} and witness w_{cp} using the same public claims, parameters, transcript-binding material, and folded-output boundary. Since the folding transition and opening material are valid, the CP relation is satisfied:

$$\mathcal{R}_{\text{CPFold}}(I_{\text{cp}}; w_{\text{cp}}) = 1.$$

The prover then encodes this CP relation as the WHIR backend statement stmt_{cp} and produces

$$\Pi_{\text{cp}} \leftarrow \text{WHIR.Prove}(\text{stmt}_{\text{cp}}, w_{\text{cp}}).$$

By completeness of the WHIR backend,

$$\text{WHIR.Verify}(\text{stmt}_{\text{cp}}, \Pi_{\text{cp}}) = 1.$$

Finally, the public accumulation verifier recomputes the public boundary, checks that stmt_{cp} is the prescribed encoding of I_{cp} , and runs the WHIR backend verifier. All checks accept, so

$$\text{AccVerify}(\text{Acc}_{\text{old}}, \{\text{Claim}_i\}_{i=1}^k, \text{Acc}_{\text{new}}, \Pi_{\text{acc}}) = 1.$$

□

7 Soundness as a Reduction to Row-Level Faithfulness

We now state the conditional soundness argument for the composed accumulator. The purpose of this section is not to prove, row by row, that the current K6a and N8 implementations satisfy the abstract accumulator relation. Instead, the goal is to reduce public soundness of the accumulated verifier to a small set of explicit proof obligations.

The most important of these obligations is typed-row faithfulness: the concrete rows encoded into the backend WHIR statement must imply the abstract folding relation of Section 5. Lemma 5.2 isolates six sufficient row-level conditions for this to hold. In this report we do not prove those six conditions for every K6a/N8 row family. Accordingly, the main theorem below should be read as a reduction theorem:

$$\begin{aligned} &\text{backend soundness} + \text{binding} + \text{transcript binding} + \text{row-level faithfulness} \\ &\implies \text{one-step accumulator soundness.} \end{aligned}$$

This distinction is important. The theorem proves that the composition is sound provided the concrete row system faithfully implements the abstract relation. It does not by itself complete the implementation-specific row audit.

Why this is not circular. The theorem does not assume the conclusion that the concrete proof is sound. It assumes a separately checkable row-level implication: accepting concrete rows must denote the same abstract WHIR-backed obligations used by $\mathcal{R}_{\text{Fold}}$. If the rows merely carry digests or similarly named fields without enforcing the corresponding verifier semantics, typed-batch faithfulness is false and the theorem cannot be applied. What remains to audit is therefore concrete and row-local: for each K6a/N8 row family, show that the public metadata, transcript roots, challenge digests, openings, folding equations, and folded-output boundary jointly imply the six obligations of Lemma 5.2.

7.1 Assumptions

We state the assumptions used in the soundness theorem.

Assumption 7.1 (WHIR backend soundness). For every efficient adversary producing a backend statement stmt_{cp} and proof Π_{cp} , if

$$\text{WHIR.Verify}(\text{stmt}_{\text{cp}}, \Pi_{\text{cp}}) = 1,$$

then, except with probability $\varepsilon_{\text{WHIR}}$, the relation encoded by stmt_{cp} is satisfiable. Equivalently, there exists a witness w_{enc} such that

$$\mathcal{R}_{\text{enc}}(\text{stmt}_{\text{cp}}; w_{\text{enc}}) = 1.$$

Assumption 7.2 (Backend-to-CP faithfulness). For every public CP instance I_{cp} , let

$$\text{stmt}_{\text{cp}} = \text{Enc}_{\text{WHIR}}(\mathcal{R}_{\text{CPFold}}, I_{\text{cp}}, \text{PP}_{\text{WHIR}}).$$

The relation represented by stmt_{cp} is faithful to $\mathcal{R}_{\text{CPFold}}$ in the following sense:

$$\exists w_{\text{enc}} : \mathcal{R}_{\text{enc}}(\text{stmt}_{\text{cp}}; w_{\text{enc}}) = 1 \implies \exists w_{\text{cp}} : \mathcal{R}_{\text{CPFold}}(I_{\text{cp}}; w_{\text{cp}}) = 1.$$

This is not an additional cryptographic assumption; it is the soundness direction of the encoding-faithfulness proof obligation from Section 5.5.

Assumption 7.3 (Commitment binding). Except with probability $\varepsilon_{\text{bind}}$, the commitments used by the folding and CP layers cannot be opened to two inconsistent values. In particular, the same public commitment cannot be used to justify two different folding witnesses, claim encodings, or accumulator states.

Assumption 7.4 (Transcript binding). Except with probability ε_{FS} , the Fiat–Shamir and folding challenges used in the proof are uniquely determined by the public transcript material checked by the verifier. Thus an accepting proof cannot use challenges derived from one public statement and be verified against another.

Assumption 7.5 (Folding soundness). Let Acc_{old} satisfy $\text{ValidAcc}(\text{Acc}_{\text{old}})$. This is an assumption about the abstract folding relation $\mathcal{R}_{\text{Fold}}$ from Section 5, applied to an abstract family of WHIR-backed claims $\{\text{Claim}_i\}_{i=1}^k$. Except with probability $\varepsilon_{\text{fold}}$, if

$$\mathcal{R}_{\text{Fold}}(\text{Acc}_{\text{old}}, \{\text{Claim}_i\}_{i=1}^k, \text{Acc}_{\text{new}}; w_{\text{fold}}) = 1,$$

then the represented input claims are valid for the declared accumulation shape and the new accumulator satisfies

$$\text{ValidAcc}(\text{Acc}_{\text{new}}).$$

For concrete K6a and N8 typed batches, identifying the public rows with those abstract claims is exactly the typed-batch faithfulness obligation, not a separate consequence of this folding assumption.

Assumption 7.6 (Typed-batch faithfulness). The concrete typed batch used by the implemented route is faithful in the sense of Definition 5.1.

7.2 Soundness Chain

The theorem below is therefore best understood as a soundness reduction. Its conclusion is only as strong as the row-level faithfulness obligation supplied for the concrete route.

Theorem 7.7 (Conditional soundness reduction for one accumulation step). *Assume Assumptions 7.1–7.6. Let B be the concrete typed batch used by an implemented route, and let*

$$\{\text{Claim}_i\}_{i=1}^k := \text{RouteMap}(B).$$

Assume $\text{ValidAcc}(\text{Acc}_{\text{old}})$. If the public accumulation verifier accepts

$$\text{AccVerify}(\text{Acc}_{\text{old}}, B, \text{Acc}_{\text{new}}, \Pi_{\text{acc}}) = 1,$$

then, except with error

$$\varepsilon_{\text{total}} \leq \varepsilon_{\text{WHIR}} + \varepsilon_{\text{fold}} + \varepsilon_{\text{bind}} + \varepsilon_{\text{FS}},$$

the following hold:

1. the represented WHIR-backed proof-validity claims $\{\text{Claim}_i\}_{i=1}^k$ are valid for the declared route shape;
2. $\text{ValidAcc}(\text{Acc}_{\text{new}})$;
3. the transition from Acc_{old} to Acc_{new} is a valid accumulator update for the represented claims.

Proof. Since the public verifier accepts, it accepts the public CP instance and checks that the backend statement is the prescribed encoding

$$\text{stmt}_{\text{cp}} = \text{Enc}_{\text{WHIR}}(\mathcal{R}_{\text{CPFold}}, I_{\text{cp}}, \text{pp}_{\text{WHIR}}).$$

It also accepts the WHIR backend proof:

$$\text{WHIR.Verify}(\text{stmt}_{\text{cp}}, \Pi_{\text{cp}}) = 1.$$

By Assumption 7.1, except with error $\varepsilon_{\text{WHIR}}$, the encoded backend relation is satisfiable. By backend-to-CP faithfulness, which is the soundness direction of the encoding-faithfulness obligation, there exists w_{cp} such that

$$\mathcal{R}_{\text{CPFold}}(I_{\text{cp}}; w_{\text{cp}}) = 1.$$

By the definition of $\mathcal{R}_{\text{CPFold}}$, and by Assumption 7.3, the witness w_{cp} contains folding witness material w_{fold} that opens the committed data consistently and satisfies

$$\mathcal{R}_{\text{Fold}}(\text{Acc}_{\text{old}}, \{\text{Claim}_i\}_{i=1}^k, \text{Acc}_{\text{new}}; w_{\text{fold}}) = 1$$

for the same public claims, transcript material, parameters, and folded-output boundary.

By Assumption 7.4, the challenges used in this folding relation are derived from the same public transcript material checked by the verifier. Hence the prover cannot prove a transition for one public transcript and present it under another.

Since $\text{ValidAcc}(\text{Acc}_{\text{old}})$, Assumption 7.5 implies that the represented claims are valid for the declared accumulation shape and that $\text{ValidAcc}(\text{Acc}_{\text{new}})$, except with error $\varepsilon_{\text{fold}}$.

Finally, Assumption 7.6 identifies the concrete typed batch B with the abstract WHIR-backed proof-validity claims $\text{RouteMap}(B)$. Therefore the accepting concrete route proof implies validity of the represented WHIR-backed claims and validity of the accumulator transition. A union bound over the backend, folding, binding, and transcript failure events gives the claimed error bound. $\square$

7.3 Error Composition

The theorem above treats faithfulness of the encoding as a proof obligation rather than as a probabilistic error. If the encoding is not faithful, then the WHIR backend may be sound for the wrong relation, and the composition theorem does not apply.

The probabilistic error terms are inherited from the component systems:

- $\varepsilon_{\text{WHIR}}$ is the soundness error of the WHIR backend for the encoded CP relation;
- $\varepsilon_{\text{fold}}$ is the soundness error of the folding reduction or accumulator update relation;
- $\varepsilon_{\text{bind}}$ captures commitment-binding failures and digest collisions that would allow inconsistent openings or public encodings;
- ε_{FS} captures Fiat–Shamir or random-oracle transcript-binding failures.

For a single accumulation step, a union bound gives

$$\varepsilon_{\text{total}} \leq \varepsilon_{\text{WHIR}} + \varepsilon_{\text{fold}} + \varepsilon_{\text{bind}} + \varepsilon_{\text{FS}}.$$

For multiple accumulation steps, the bound scales with the number of steps unless a tighter accumulation-specific analysis is used.

Long-chain accumulation. The theorem above is a one-step statement. For a chain of T accumulation steps, let

$$\text{Acc}_0 \rightarrow \text{Acc}_1 \rightarrow \cdots \rightarrow \text{Acc}_T$$

be the public accumulator sequence, and let B_t be the typed batch absorbed at step t . A direct application of Theorem 7.7 to each step gives

$$\varepsilon_{\text{chain}} \leq T \cdot (\varepsilon_{\text{WHIR}} + \varepsilon_{\text{fold}} + \varepsilon_{\text{bind}} + \varepsilon_{\text{FS}}).$$

This is the conservative bound used by the present report.

A tighter analysis would treat the entire chain as one composed experiment rather than as T unrelated applications of the one-step theorem. Such an analysis would maintain an invariant

$$\text{ValidAcc}(\text{Acc}_t)$$

for every step t , and would prove that an accepting transition preserves this invariant except on a global bad event. The goal would be to charge failure not once per step, but once per underlying source of failure.

There are three natural places where such a refinement could improve the bound.

First, for algebraic folding errors, one may try to bound the probability that any invalid folding transition passes by using the aggregate degree of the polynomial identities checked across the whole chain. In favorable settings this can yield a bound of the form

$$\frac{\sum_{t=1}^T d_t}{|\mathbb{F}|}$$

rather than T times a worst-case per-step estimate, where d_t is the relevant degree contribution of step t .

Second, for random-oracle and Fiat–Shamir failures, one can analyze the global transcript as a single domain-separated random-oracle experiment. If every step binds the previous accumulator digest, the step index, the batch digest, and the route identifier into the challenge derivation, then a successful transcript mismatch must correspond to either a digest collision or an inconsistency in the global transcript tree. This suggests a bound in terms of the total number of random-oracle queries and digest outputs, rather than a black-box multiplication by T .

Third, for commitment binding, the relevant bad event is that some commitment in the entire chain admits two inconsistent openings. If the commitments are globally indexed by accumulator

step and route domain, then the binding loss can be bounded over the total number of commitments appearing in the chain.

The present implementation does not require this tighter analysis, because the evaluation studies bounded one-step accumulation routes. Nevertheless, a production long-chain accumulator should prove a chain-level invariant theorem: if $\text{ValidAcc}(\text{Acc}_0)$ holds and all T public transitions verify, then $\text{ValidAcc}(\text{Acc}_T)$ holds and all represented claims in the chain are valid, except with a global error bound derived from the full transcript and commitment experiment.

8 Implementation

This section maps the abstract construction of Sections 4 and 5 to the current codebase. The implementation does not expose a generic accumulator over arbitrary vectors of WHIR proof objects. It exposes typed public batches and route-specific accumulator instances whose rows are intended to represent the WHIR-backed obligations of the abstract construction.

8.1 Routes and Data Model

There are three relevant implementation routes. The default `prove_public/verify_public` route is the product WHIR+WHIR public verifier. K6a is the explicit single-oracle full-workload accumulator route. N8 is the explicit integrated multi-oracle alternative with ACC.P, ACC.V, and ACC.D entry points. K6a and N8 are opt-in NonZK integrity routes; they do not silently replace the default product verifier.

The implementation represents the abstract objects

$$\text{Claim}_i, \quad \text{Acc}_{\text{old}}, \quad \text{Acc}_{\text{new}}, \quad I_{\text{cp}}, \quad \Pi_{\text{cp}}$$

by typed public structures rather than by raw tuples:

Abstract object	Implementation role
Claim_i	typed WHIR-backed obligation represented in the batch
Acc_{old}	old accumulator public instance
Acc_{new}	new accumulator public instance
I_{cp}	public instance for the typed folding/CP relation
w_{cp}	private folding witness, openings, and auxiliary data
Π_{cp}	WHIR backend proof for the encoded CP relation

Route-specific batches encode the public obligations corresponding to the abstract WHIR statement/proof claims, plus the metadata needed to bind them to an accumulator transition. This is why typed-batch faithfulness is the main implementation obligation: the typed public rows must denote the same public verification obligations as the abstract claim notation, not just carry digests with compatible names.

8.2 Single-Oracle versus Multi-Oracle Organization

The terms *single-oracle* and *multi-oracle* refer to the organization of the backend row system, not to whether the public route is allowed to verify many independent proofs.

In K6a, the accumulation workload is packaged into one public-canonical backend table and one top-level WHIR backend proof for the typed CP relation. In N8, the same public transition is organized into logical row families before being proved by one integrated WHIR proof. These families include K6a semantic rows, tuple-RLC semantic rows, and transition/binding rows.

Thus the distinction is not

one proof versus many independent proofs.

It is

one monolithic row/table organization

versus

multiple logical oracle families inside one integrated route.

N8 is therefore evidence for the native multi-oracle direction because it structures the backend statement without turning the route into a list of independent WHIR verifications.

8.3 Encoded Folding Transition

The concrete routes encode the abstract transition

$$\mathcal{R}_{\text{Fold}}(\text{Acc}_{\text{old}}, \{\text{Claim}_i\}_{i=1}^k, \text{Acc}_{\text{new}}; w_{\text{fold}}) = 1.$$

The typed rows must enforce shape compatibility, challenge derivation from the declared public transcript, opening consistency, algebraic folding consistency, and equality between the computed folded state and Acc_{new} . The CP relation is the public-verifier form of this transition: it proves those checks while keeping folding witnesses, openings, and auxiliary prover data outside the verifier boundary.

WHIR is then used as the backend proof system for the encoded CP relation. The implementation forms

$$\text{stmt}_{\text{cp}} = \text{Enc}_{\text{WHIR}}(\mathcal{R}_{\text{CPFold}}, I_{\text{cp}}, \text{pp}_{\text{WHIR}})$$

and verifies Π_{cp} against the recomputed statement. The backend statement must bind the same public data as the CP instance: old and new accumulators, typed batch, relation metadata, transcript roots, challenge digests, and folded-output boundary. Omitting any of these values would make the backend proof sound for a weaker statement than the advertised accumulator transition.

8.4 Tests and Guardrails

The main implementation guardrail is fail-closed public verification. A public route should reject if any required backend authority, typed verification, or public transcript binding is unavailable. It should not fall back to witness-side replay, debug helpers, or non-authoritative typed checks.

The corresponding review matrix is:

1. modifying public inputs or relation metadata invalidates the proof;
2. modifying transcript roots or challenge digests invalidates the proof;
3. modifying the old or new accumulator invalidates the proof;
4. using a backend statement inconsistent with the CP instance is rejected;
5. missing typed backend support causes rejection rather than fallback;
6. explicit accumulation routes remain opt-in and do not silently replace the default public verifier.

These are the implementation counterparts of the binding assumptions used in the soundness theorem.

9 Evaluation

This section evaluates the implemented WHIR accumulation routes against the default public WHIR verification baseline. The goal is to distinguish the cost of the authoritative public boundary from the cost of the explicit accumulator routes. The central empirical question is whether the accumulator routes behave like one public transition, rather than like repeated monolithic public verification.

We evaluate three implementation routes:

1. the default WHIR+WHIR product public verifier, the authoritative baseline exposed by `verify_public` and `verify_v2`;
2. the explicit K6a SYMBT3 accumulator route, which is the current full-workload NonZK integrity accumulation path;
3. the explicit N8 integrated route and native multi-oracle instrumentation, which show the current one-WHIR and native tuple-leaf direction.

The evaluation reports five quantities:

prover time, verifier time, proof size, public statement size, scaling with arity.

We also report a qualitative bottleneck analysis explaining which part of the construction dominates the current cost profile.

9.1 Experimental Setup

All experiments in this section were run locally on May 25, 2026. The checkout was dirty before the benchmark run, so the commit hash should be read as a code base identifier rather than a clean-release tag.

Parameter	Value
Machine	MacBook Pro, Apple M3 Pro, 12 cores
Memory	36 GB
Operating system	macOS 26.3.1, Darwin 25.3.0, arm64
Rust toolchain	rustc 1.93.1, cargo 1.93.1
Repository commit	477ee55, dirty worktree
Backend / field	WHIR over BabyBear; Poseidon2/BabyBear public digests
Bench arities	$k \in \{1, 2, 4, 8, 16, 32, 64\}$ for K6a/N8; product comparison through $k = 8$
Measurement mode	Criterion wall-clock benchmarks plus profiled CSV rows

Table 2: Local benchmark environment.

The exact commands were:

```
cargo bench --bench whir_scaling --features whir --no-run
SYMPHONY_WHIR_PUBLIC_VERIFY_KS=1,2,4,8,16,32,64 cargo bench --bench whir_scaling --features whir -- "
  symbt3_accumulator_authority_vs_k"
SYMPHONY_WHIR_PUBLIC_VERIFY_KS=1,2,4,8,16,32,64 cargo bench --bench whir_scaling --features whir -- "
  symbt3_n8_integrated_authority_vs_k"
SYMPHONY_WHIR_NATIVE_ORACLE_COUNTS=1,2,4,8,16,32,64 cargo bench --bench whir_scaling --features whir -- "
  instrumented_multi_oracle"
SYMPHONY_WHIR_PUBLIC_VERIFY_KS=1,2,4,8,16,32,64 cargo bench --bench whir_scaling --features whir -- "
  product_route_comparison_vs_k"
SYMPHONY_WHIR_PUBLIC_VERIFY_KS=16 cargo bench --bench whir_scaling --features whir -- "product_route_comparison_vs_k"
python3 scripts/plot_report_figures.py
```

The product-comparison command emitted rows through $k = 8$ and was then killed by the operating system at $k = 16$; an isolated $k = 16$ retry was also killed. The local logs did not record a final resource counter for the killed $k = 16$ product run, so we treat this row as right-censored data rather than extrapolating a product baseline beyond $k = 8$. The most likely cause is memory pressure on the 36 GB machine from the large monolithic typed-CP row system, but the local run did not record an OOM or watchdog counter, so this remains a hypothesis rather than a measured diagnosis. The product table therefore reports the completed product rows only. The K6a and N8 accumulator routes completed the full requested k -set.

Measurement protocol. The benchmark rows were produced by Criterion-style wall-clock benchmarks and CSV instrumentation. Where Criterion logs are available, the benchmarks use a 3 second warmup and collect 10 samples per arity. For example, the N8 integrated authority run records 10 samples for each measured arity $k \in \{1, 2, 4, 8, 16, 32, 64\}$. The tables report the summarized wall-clock values emitted by the benchmark scripts.

The product-route comparison CSV does not currently include per-sample standard deviations or confidence intervals. Therefore, the evaluation should be read as a route-comparison measurement rather than as a statistically complete microbenchmark study. The large product-versus-K6a verifier gap is robust at the scale reported in Table 3, but smaller differences, such as K6a versus N8 at nearby arities, should not be over-interpreted without repeated-run variance data.

The benchmark artifacts used below are:

- benchmarks/product_route_comparison.csv;
- benchmarks/symbt3_scaling.csv;
- benchmarks/symbt3_instrumented_benchmark.jsonl;
- benchmarks/n8_integrated_authority.csv;
- benchmarks/n8_integrated_timer.csv;
- benchmarks/symbt3_instrumented_multi_oracle.jsonl.

9.2 Product Public-Verifier Baseline

The default product route is the authoritative WHIR+WHIR public verifier. It uses public data only and is expected to pass for honest WHIR public proofs and fail closed for malformed public data. Its cost is the right baseline because it is the product public-verifier boundary, even though it is not the final accumulator route.

Table 3 compares this default product route with the explicit K6a accumulator route for the product arities that completed locally. The product route is expensive: verifier time is about 1.9 seconds at $k = 1$ and 30.9 seconds at $k = 8$ on this machine. This reflects the current monolithic typed-CP public verifier, not witness-side replay. The attempted $k = 16$ product baseline did not complete locally because the benchmark process was killed by the operating system.

k	Product verify ms	K6a verify ms	Verify speedup	Product proof bytes	K6a proof bytes	Product public bytes	K6a public bytes
1	1,896.276	17.300	109.6x	1,206,359	318,649	15,171	18,715
2	4,024.325	19.101	210.7x	1,257,211	331,889	15,187	18,715
4	9,658.127	21.988	439.2x	1,556,382	319,119	15,219	18,715
8	30,894.046	29.518	1,046.6x	1,612,774	376,078	15,283	18,715

Table 3: Default product public verifier compared with the explicit K6a NonZK integrity accumulator route. Times and sizes are from benchmarks/product_route_comparison.csv.




Figure 2: Verifier time versus accumulated arity. K6a and N8 nearly overlap in this benchmark; both are shown with distinct markers to make clear that they are separate explicit routes. The product public verifier is the authoritative baseline and only completed through $k = 8$ locally.

9.3 K6a Explicit Accumulator Route

The K6a route replaces the monolithic product public-verifier workload by one explicit accumulator transition:

$$\text{AccVerify}(\text{Acc}_{\text{old}}, \{\text{Claim}_i\}_{i=1}^k, \text{Acc}_{\text{new}}, \Pi_{\text{acc}}) = 1.$$

This route is explicit opt-in, NonZK integrity-only, and not the default `verify_public` route. In the benchmarked shape it verifies successfully for all measured arities, selects the product route, and does not use the monolithic fallback.

The *Source claims* column in Table 4 is the logical source R1CS residual-claim count, not the number of top-level WHIR proofs. In this workload each accumulated item contributes 64 such logical source claims, so the counter is $64k$; the verifier-side source residual evaluation is compressed to one batched evaluation in the current K6a profile.




Figure 3: Prover time versus accumulated arity for the product baseline, K6a accumulator, and N8 integrated route. Product-route rows stop at $k = 8$ because the local $k = 16$ product benchmark was killed by the operating system.

k	Prove ms	Verify ms	Proof bytes	Public bytes	Oracle len.	Source claims
1	16.538	16.161	307,597	18,715	4,096	64
2	19.493	18.981	345,133	18,715	4,096	128
4	24.232	22.168	340,214	18,715	4,096	256
8	34.892	29.320	372,764	18,715	8,192	512
16	54.309	43.982	421,516	18,715	16,384	1,024
32	95.278	72.256	645,435	18,715	32,768	2,048
64	174.264	125.655	682,806	18,715	65,536	4,096

Table 4: K6a explicit accumulator scaling from benchmarks/symt3_scaling.csv. The public statement size is flat for this run, while the logical source-claim counter scales with k .

The generated scaling summary reports log-log slopes of 0.490 for verification time, 0.569 for proving time, 0.199 for proof bytes, and 0.000 for public statement bytes. The architecture guardrail also remains stable: one top-level WHIR proof, zero family-columnar subproofs, and one backend table for every measured k .

9.4 N8 Integrated Route and Native Multi-Oracle Evidence

The N8 route is an explicit integrated one-WHIR alternative for same-shape NonZK accumulation transitions. It is not the default product `verify_public` route, but it is useful evidence for the direction of the implementation: the public route can be expressed through a more integrated native-oracle structure rather than only through the K6a wrapper-style path.




Figure 4: Serialized proof size versus arity. The accumulator routes keep a single public transition proof shape rather than a list of independent WHIR proofs.




Figure 5: Public statement size versus arity. K6a and N8 are flat in the measured shape because the accumulator boundary is fixed-size for this workload.

k	Status	Prove ms	Verify ms	Proof bytes	WHIR inst.	Roots	Soundness bits
1	OK	24.979	14.670	640,702	1	1	124
2	OK	32.787	17.446	686,949	1	1	124
4	OK	50.262	20.378	726,636	1	1	124
8	OK	81.866	29.977	956,339	1	1	124
16	OK	151.373	44.832	1,001,709	1	1	124
32	OK	280.310	69.956	1,044,400	1	1	124
64	OK	558.118	122.412	1,091,966	1	1	124

Table 5: N8 integrated explicit route from benchmarks/n8_integrated_authority.csv. All measured rows pass the N8 authority gate with complete K6a, tuple-RLC, and transition semantics.




Figure 6: N8 timing breakdown for representative arities. The dominant components identify the remaining implementation bottlenecks.

The native multi-oracle instrumentation is not the main accumulator workload, but it records the intended tuple-leaf shape. In the refreshed JSONL run over logical-oracle counts 1, 2, 4, 8, 16, 32, 64, every true native tuple-leaf row passes the shape guard and uses one WHIR instance, one root, one query schedule, and one transcript. For the $k_{\text{table}} = 8, 64$ -logical-oracle row, the tuple-leaf verifier measurement is 0.516 ms against 30.683 ms for the corresponding n -times-single-oracle baseline.

9.5 Bottleneck Analysis

The current implementation identifies three main bottlenecks.

First, the authoritative product public verifier remains expensive because the monolithic typed-CP relation is large. This route is still the correct public baseline: a first unoptimized, monolithic implementation that is expected to pass for honest public WHIR proofs and fail closed for malformed public data. The local timings nevertheless show seconds of verifier work even for small k . The local monolithic run did not reach $k = 16$, which reinforces that this baseline is not the intended accumulated cost profile.

Second, public binding material is not free. In the K6a route the public statement is 18,715 bytes for all measured k , slightly larger than the product public envelope. The main public-byte contributors are folded values, the accumulator boundary, and relation/output digests. This material is soundness-critical and should not be removed without an equivalence audit.

Third, the explicit accumulator routes are NonZK integrity-only. K6a is the current full-workload accumulation path and N8 is the integrated explicit alternative, but neither should be described as the default product `verify_public` route or as a production zero-knowledge accumulator.

The evaluation should therefore be read as a measurement of current route shape and public-verifier cost, not as the final performance profile of an optimized native multi-oracle accumulator. The key result is nevertheless clear: the explicit accumulator routes check one public transition and avoid the repeated monolithic public-verifier cost that dominates the baseline.

10 Challenges and Limitations

The construction and soundness argument above are conditional. They identify the interface required for a sound WHIR-backed accumulator, but they do not by themselves certify every concrete route as production-authoritative. The main limitations are the following.

10.1 Relation and Transcript Faithfulness

The central limitation is faithfulness of the encoded relation. The WHIR backend is sound for the statement it actually proves; if that statement omits an accumulator field, transcript digest, opening condition, or folded-output boundary, then it may prove a weaker language than the accumulator API claims. The K6a and N8 rows must therefore be checked against the ledger of Section 5: claim representation, shape consistency, challenge binding, opening consistency, folding consistency, and output consistency.

This includes commitment compatibility. Lattice/Ajtai commitments, CP witness commitments, and WHIR oracle commitments bind different objects. Soundness requires explicit public digests and relation identifiers tying all of them to the same transition. It also includes transcript binding: folding challenges, WHIR backend challenges, and public verifier digests must be derived from the same public statement material, and the verifier must reject if this cannot be checked from public data.

10.2 Prototype Boundaries

The abstract notation uses a batch of WHIR proof-validity claims

$$\{\text{Claim}_i\}_{i=1}^k,$$

but the concrete single-oracle and multi-oracle routes consume typed accumulation batches. They are implementations of the semantic interface, not unrestricted generic accumulators for arbitrary WHIR proof objects. The current K6a and N8 routes should therefore be viewed as soundness-guided NonZK integrity prototypes, not as a complete zero-knowledge Symphony-style SNARK.

10.3 Remaining Gaps

The main remaining gaps are the following.

1. The typed CP relation remains large, which motivates the native multi-oracle direction measured in the evaluation.

2. The row-level faithfulness audit remains the main mathematical gap between the reduction theorem and concrete K6a/N8 soundness.
3. The current routes are not zero-knowledge. Adding ZK would require masking the folding witness, openings, and auxiliary CP witness data, then proving the masked relation while binding masking commitments into the same transcript.
4. The soundness theorem is one-step. Long-chain accumulation needs a `ValidAcc` invariant and a global transcript and commitment analysis rather than only the conservative union bound of Section 7.3.
5. A production-authoritative route still needs a complete public-verifier audit: no witness-side replay, no debug-only fallback, no non-authoritative helper, and explicit binding of every public statement component.

11 Related Work

This report sits at the intersection of four lines of work: WHIR-style hash-based proof backends, lattice-based folding, commit-and-prove SNARKs, and accumulation schemes. We discuss each in turn and explain how the present construction differs from prior work.

11.1 WHIR and Reed–Solomon Proximity Testing

WHIR [ACFY24b] is an interactive oracle proof of proximity for constrained Reed–Solomon codes with small query complexity and very fast verification. It can serve as a backend for polynomial-IOP style proof systems, replacing or improving on earlier Reed–Solomon proximity-testing components such as FRI [BSBHR18], STIR [ACFY24a], and BaseFold-style folding [ZCF23]. Its verifier efficiency makes it attractive as a public proof backend.

The present report uses WHIR in two different roles. First, the input objects being accumulated are WHIR-backed proof-validity claims. Second, WHIR is used again as the backend proof system for the commit-and-prove relation that certifies the folding transition. This differs from using WHIR only as a standalone proof system: the central question here is whether WHIR verification claims can themselves be accumulated without losing the public soundness boundary.

11.2 Symphony and Lattice-Based Folding

Symphony [Che25] proposes a lattice-based high-arity folding framework for building scalable SNARKs in the random oracle model. Its key architectural idea is to use high-arity folding together with a commit-and-prove compiler, while avoiding the need to place Fiat–Shamir hash computations inside the proved statements. This makes Symphony a natural starting point for studying memory-efficient and potentially post-quantum accumulation architectures.

Our construction follows the Symphony-style viewpoint that a folding transition should be certified by a CP relation, rather than by exposing prover-side folding witnesses to the public verifier. However, the goal here is narrower: we do not claim to implement a full production Symphony SNARK. Instead, we study the specific composition problem that arises when WHIR-backed proof-validity claims are folded through a Symphony-style accumulator and the resulting CP relation is itself proved with WHIR.

11.3 Commit-and-Prove SNARKs

Commit-and-prove SNARKs allow a prover to prove a relation while also proving consistency with previously committed data. In the present construction, the CP layer is the bridge between the semantic folding relation and the public backend proof. It certifies that the prover-side folding witness, openings, and auxiliary data are consistent with the public accumulator transition, while keeping those objects outside the public verifier boundary.

The CP layer is also where the main faithfulness requirement arises. The backend proof system is sound only for the relation it actually encodes. Therefore, the WHIR backend statement must encode exactly the CP folding relation for the same public instance: the same old accumulator, new accumulator, claim batch, transcript-binding material, and folded-output boundary. This faithfulness condition is what prevents the backend from proving a weaker relation than the one required by the accumulator.

11.4 Accumulation and Folding Schemes

Accumulation and folding schemes are widely used to reduce the cost of verifying many related proof obligations. Instead of recursively verifying every proof inside another proof system, an accumulator maintains a smaller state that represents the validity of many prior claims. This idea underlies many modern incrementally verifiable computation and proof-carrying-data systems.

Recent work explores different accumulation tradeoffs. WARP gives a hash-based accumulation scheme with linear prover time and logarithmic verifier time [BCFW25]. Quasar studies sublinear accumulation for multiple instances, aiming to reduce the verifier work associated with accumulating many predicate instances [ZGC⁺25]. These works are close in spirit to the present report because they also target the repeated-verification bottleneck.

The difference is the object being accumulated. This report focuses specifically on WHIR-backed proof-validity claims and on their composition with a Symphony-style lattice-folding and CP-SNARK architecture. The main question is not only whether accumulation can reduce verifier work, but whether the verification relation induced by WHIR proofs can be folded and re-proved with WHIR while preserving the public verifier boundary.

11.5 Transparent SNARKs and Spartan

Spartan [Set19] is an important transparent SNARK construction for R1CS that combines sumcheck-style protocols, polynomial commitments, and preprocessing ideas to obtain succinct verification without a trusted setup. Spartan is relevant background for this report because it illustrates a recurring theme in modern proof systems: arithmetize the computation, bind the public relation carefully, and use polynomial-commitment machinery to obtain succinct verification.

The present work is not a Spartan-style construction. It does not build a new R1CS SNARK from scratch. Instead, it assumes a WHIR-backed proving stack and studies whether the resulting proof-validity claims can be accumulated through lattice-based folding. Spartan is therefore best viewed as part of the broader context of transparent succinct proof systems, rather than as a direct component of the construction.

11.6 Positioning of This Work

The contribution of this report is not a new low-degree test, a new polynomial commitment scheme, or a complete new SNARK. Its contribution is a soundness- and implementation-oriented study of a

particular composition:

WHIR proof-validity claims $\longrightarrow$ lattice-based folding
 $\longrightarrow$ commit-and-prove relation
 $\longrightarrow$ WHIR backend proof.

The main novelty is the interface question. WHIR and Symphony each provide useful components, but their composition is not automatic. The CP relation must prove the correct folding transition, the WHIR backend must faithfully encode that CP relation, and the public verifier must remain fail-closed and authoritative. This report makes those obligations explicit and evaluates the current implementation routes against them.

12 Future Work

The implementation and analysis in this report suggest several directions for future work.

Faithfulness proof for typed rows. The soundness theorem in this report assumes that the typed SYMBT3/K6a/N8 rows faithfully encode the abstract CP folding relation. Section 5.7 reduces this to six row-level obligations: claim representation, shape consistency, challenge binding, opening consistency, algebraic folding consistency, and output consistency. A natural next step is to prove these obligations directly for the implemented rows. Such a proof should proceed row family by row family, showing that every public binding condition in $\mathcal{R}_{\text{CPFold}}$ is represented in the backend WHIR statement. This would close the main gap between the abstract soundness theorem and the concrete implementation.

Quasar-style partial evaluation. Quasar is an especially interesting exploratory direction for this project. The current accumulator routes still carry a large typed CP relation, and the evaluation shows that this relation remains a dominant cost. Quasar-style multi-instance accumulation suggests a different way to reduce verifier work: replace some random-linear-combination style verifier operations with partial evaluation techniques that are sublinear in the number of accumulated instances. An interesting next step would be to study whether the WHIR-backed CP relation can be reorganized so that the accumulator verifier checks a smaller partially evaluated object rather than a monolithic folded relation. This would not replace the soundness interface developed in this report, but it could provide a more efficient realization of the same public-boundary requirements.

Zero-knowledge. The current routes are NonZK integrity-only. Adding zero-knowledge would require masking or blinding the relevant witness material, proving that the resulting distribution hides the private folding data, and checking that the WHIR backend and transcript-binding logic remain compatible with the zero-knowledge layer. This would also require revisiting the CP relation, since the current report focuses on public soundness and accumulator integrity rather than witness privacy.

Long-chain accumulation. The current soundness theorem is a one-step theorem. Section 7.3 explains the conservative union-bound extension to T steps and sketches how a sharper chain-level analysis might be obtained. A full treatment should define a chain invariant $\text{ValidAcc}(\text{Acc}_t)$, prove that each accepted transition preserves it, and analyze the entire random-oracle, commitment, and folding experiment globally rather than as T independent one-step experiments.

Broader comparisons. The evaluation in this report compares the product public verifier, the K6a explicit accumulator route, and the N8 integrated route. A broader comparison would benchmark the same workloads against other accumulation strategies, including WARP-style hash-based accumulation and Quasar-style multi-instance accumulation. Such a comparison would clarify whether the main cost in this project comes from WHIR backend proving, typed CP relation size, transcript binding, or the specific lattice-folding interface.

13 Conclusion

This report studied whether WHIR proof-validity claims can be accumulated through Symphony-style lattice-based folding while preserving a public, fail-closed verifier boundary. The central object is the WHIR verification claim, including the public transcript and folded-output data that make it meaningful.

We formalized this composition as a chain of relations: WHIR verification claims are absorbed by a folding relation, the folding transition is certified by a commit-and-prove relation, and that CP relation is proved using WHIR as the backend. We proved completeness and gave a conditional soundness theorem under backend soundness, folding soundness, commitment binding, transcript binding, typed-batch faithfulness, and faithful encoding of the CP relation into WHIR.

The implementation results show that the public-verifier boundary can be preserved in explicit accumulation routes. In the measured workloads, the K6a route avoids the repeated monolithic public-verifier cost of the product baseline, while N8 demonstrates an integrated one-WHIR route with stable authority metadata across arities. At the same time, the evaluation confirms that the current system is still a prototype: the typed CP relation, transcript binding, and relation encoding remain the main bottlenecks toward a production-authoritative WHIR-backed accumulator.

References

- [ACFY24a] Gal Arnon, Alessandro Chiesa, Giacomo Fenzi, and Eylon Yogev. STIR: Reed–Solomon Proximity Testing with Fewer Queries. Cryptology ePrint Archive, Paper 2024/390, 2024.
- [ACFY24b] Gal Arnon, Alessandro Chiesa, Giacomo Fenzi, and Eylon Yogev. WHIR: Reed–Solomon Proximity Testing with Super-Fast Verification. Cryptology ePrint Archive, Paper 2024/1586, 2024.
- [BCFW25] Benedikt Bünz, Alessandro Chiesa, Giacomo Fenzi, and William Wang. Linear-time accumulation schemes. Cryptology ePrint Archive, Paper 2025/753, 2025.
- [BSBHR18] Eli Ben-Sasson, Iddo Bentov, Yinon Horesh, and Michael Riabzev. Fast Reed–Solomon Interactive Oracle Proofs of Proximity. In *45th International Colloquium on Automata, Languages, and Programming (ICALP 2018)*, volume 107 of *Leibniz International Proceedings in Informatics (LIPIcs)*, pages 14:1–14:17. Schloss Dagstuhl – Leibniz-Zentrum für Informatik, 2018.
- [Che25] Binyi Chen. Symphony: Scalable SNARKs in the Random Oracle Model from Lattice-Based High-Arity Folding. Cryptology ePrint Archive, Paper 2025/1905, 2025.
- [Set19] Srinath Setty. Spartan: Efficient and general-purpose zkSNARKs without trusted setup. Cryptology ePrint Archive, Paper 2019/550, 2019.

- [ZCF23] Hadas Zeilberger, Binyi Chen, and Ben Fisch. BaseFold: Efficient Field-Agnostic Polynomial Commitment Schemes from Foldable Codes. Cryptology ePrint Archive, Paper 2023/1705, 2023.
- [ZGC⁺25] Tianyu Zheng, Shang Gao, Sherman S. M. Chow, Yu Guo, and Bin Xiao. Quasar: Sublinear multi-cast commitment mixing in recursive accumulation. Cryptology ePrint Archive, Paper 2025/1912, 2025.